

! This class has been made inactive. No posts will be allowed until an instructor reactivates the class.

note @77

stop following 65 views

Actions

Lab 8 Questions

Hi all,

Please post to this discussion for Lab 8 clarifications.

labs

Edit good note | 0

Updated 12 months ago by Shurui Zhou

followup discussions, for lingering questions and

2 endorsed followup comments

comments

☒ Resolved ☐ Unresolved @77_f1

Actions

(Anon. Beaker to classmates) 12 months ago

Hi, I can't read the content following this link: <https://medium.com/swlh/deploy-a-machine-learning-model-with-aws-elasticbeanstalk-dfcc47b6043e> for free.

You've read all of your member-only stories this month. Become a member to read and support the writers and publications uncovering new insights in the topics that matter to you.

Upgrade

helpful! | 0

12 months ago

Actions

Have you tried Incognito mode? Use a different device? Use a different IP address (U of T network, home network, Tim Hortons Wi-Fi, etc.)? I'd be surprised if you could not outsmart them some way.

helpful! | 2

months ago

Actions

Yes, going incognito will work.

good comment | 0

Reply to this followup discussion

☒ Resolved ☐ Unresolved @77_f2

Actions

[REDACTED] (Anon. Beaker to classmates) 12 months ago

Can anyone explain more about how to use the predefined model? I'm able to run their model and get the output. But the output doesn't tell info about the final label that whether the input should be labeled as 1 or 0.

For this project, you will not have to train your own model. You may choose to load this Fake News detector model from the Hugging Face Model library: <https://huggingface.co/mrm8488/bert-tiny-finetuned-fake-news-detection>, this page includes a code snippet that you can use to load the model in your environment.

How to use from the [Hugging Face](#) /transformers library

 Copy to clipboard

```
from transformers import AutoTokenizer, AutoModelForSequenceClassification

tokenizer = AutoTokenizer.from_pretrained("mrm8488/bert-tiny-finetuned-fake-news-detection")

model = AutoModelForSequenceClassification.from_pretrained("mrm8488/bert-tiny-finetuned-fake-news-detection")
```

helpful! | 1

[REDACTED] (classmates) 12 months ago

Actions

+1 same question

helpful! | 0

[REDACTED] months ago

Actions

The output gives you output probabilities, so just take the maximum probability as the label.

good comment | 0

[REDACTED] 2 months ago

Actions

How do you get the output from the model? I ran the tokenizer on the text with return_tensors="pt" and then ran the model on the tokenized input. I get "SequenceClassifierOutput(loss=None, logits=tensor([[-3.3401, 3.8496]], grad_fn=<AddmmBackward0>), hidden_states=None, attentions=None)" using the example you've given below but how do I interpret this?

helpful! | 0

[REDACTED] months ago

Actions

Read through this, it tells you how to interpret the output:

https://huggingface.co/transformers/main_classes/output.html#transformers.modeling_outputs.SequenceClassifierOutput

good comment | 0

[REDACTED] months ago

Actions

Don't worry about using huggingface, I am updating the assignment to no longer require this.

good comment | 0

Reply to this followup discussion

☒ Resolved ☐ Unresolved @77_f3

Actions

(Anon. Beaker to classmates) 12 months ago

Has anyone able to get this API(<https://huggingface.co/mrm8488/bert-tiny-finetuned-fake-news-detection>) generate a higher score for Label0? I've tried lots of messages, it keeps giving me a 0.99 score for Label1 and 0.01 score for Label0.

helpful! | 1

12 months ago

Actions

I haven't been able to either...

helpful! | 0

months ago

Actions

You are not expected to train the model nor are you evaluated on what the model outputs, so just use the model with whatever behavior it gives you.

good comment | 0

months ago

Actions

Also here is a high Label 0 example: "The Trump campaign has confirmed to Hannity.com that Mr. Trump did indeed send his plane to make two trips from North Carolina to Miami, Florida to transport over 200 Gulf War Marines back home." this is indeed fake news btw. It is okay if you find examples from online sources for your test cases. Please make sure to include whatever your source is for the examples if you do use them.

good comment | 0

12 months ago

Actions

If this example is fake news, then shouldn't it be a Label 1 then? In which case it's fake news. I think the question is if there are any inputs that can generate a higher score for Label 0 which is real news. Tried putting in real news into the API mentioned but still getting a label 1 (fake news).

helpful! | 0

months ago

Actions

I wouldn't worry about the actual data that produces the output, the model is very lightweight, we can't expect it to be very good.

good comment | 0

Reply to this followup discussion

☒ Resolved ☐ Unresolved @77_f4

Actions

(Anon. Scale to classmates) 12 months ago

Anyone try following the Medium tutorial and keep getting an error when using the zip file mention in their GitHub repository? The deployment on AWS keeps failing

helpful! | 0

months ago

Actions

The error is a MemoryError on AWS since we need to pip install torch to use AutoModelForSequenceClassification and since the torch package is >800mb, and the limit of the RAM for t3.micro is 1gb, there doesn't seem to be enough space to handle it

helpful! | 0

Actions ▾

I have checked this, and it turns out you are correct, getting PyTorch to load in t3.micro is infeasible for the scope of this assignment. I am updating the instructions on the assignment and providing a model file you can use with sci-kit learn, which should only need a few megabytes and is within the constraints to t3.micro. Stay tuned for that update.

good comment | 0

Actions ▾

The instructions have been updated. The new model should load perfectly fine in the t3.micro instance and no longer requires PyTorch. Please look at the updated instructions as per the announcement. Thanks for catching this Andy.

good comment | 0

Actions ▾

I have a different error where the initialization fails when I upload the zip file. The only errors I can find in the deployment logs (specifically, cfn-init-cmd.log) are:

```
2021-11-29 06:12:57,500 P2818 [INFO] Command 02_dev_libs
2021-11-29 06:12:58,090 P2818 [INFO] -----Command Output-----
2021-11-29 06:12:58,091 P2818 [INFO] Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
2021-11-29 06:12:58,091 P2818 [INFO] Warning: group Development Libraries does not exist.
2021-11-29 06:12:58,091 P2818 [INFO] Maybe run: yum groups mark install (see man yum)
2021-11-29 06:12:58,091 P2818 [INFO] Error: No packages in any requested group available to install or update
2021-11-29 06:12:58,091 P2818 [INFO] -----
2021-11-29 06:12:58,091 P2818 [ERROR] Exited with error code 1
```

The 02_dev_libs command in 00.application.py (sudo yum groupinstall -y "Development Libraries") is failing. I have tried a bunch of workarounds online but there is not much out there, and none of them worked. I have gone through the environment set up in the medium article multiple times trying to figure out if I did something wrong. Has anyone else run into this?

helpful! | 0

Actions ▾


When you create the zip, make sure to make the zip by selecting the files within the folder directly. If you zip the entire folder containing your code, it will not work. I.e, unzip your zip folder, and click on the unzipped folder, if there is another directory inside before you get to your python file you've zipped incorrectly.

good comment | 0

Reply to this followup discussion

☒ Resolved ☐ Unresolved @77_f5 

Actions ▾

 non. Mouse to classmates) 12 months ago

Just wondering if I got full marks on all the previous labs, does Lab 8 add bonus so that the total exceeds 15% or is it just a compensation if you lost marks on previous ones?

~ An instructor (Shurui Zhou) thinks this is a good comment ~

undo helpful 4

months ago

Actions ▾

a compensation if you lost marks on previous ones

good comment | 1

 2 months ago

Actions ▾

Thanks for asking this question. I assumed right away that the total will exceed 15% and cover lost marks in other assignments. I was wrong. Therefore, this lab will have no value to a student who already has 100% marks in the labs.

helpful! | 0

 ns ago

Actions ▾

Will Lab8 compensate for other documenting assignments or it only compensates for labs?

helpful! | 2

1 months ago

Actions ▾

This will only be compensated for labs.

good comment | 0

Reply to this followup discussion

☒ Resolved ☐ Unresolved @77_f6 

Actions ▾

Actions ▾

Reply to this followup discussion

Start a new followup discussion

Compose a new followup discussion

